



# Group 5

## Team Batak (TB)

# Table of Contents

## Meet the Team

## SageMaker Deployment Options

- ☐ Real-Time Inference Endpoint
- ☐ Batch Transform

## Amazon SageMaker Cookbook

- ☐ Chapter 1
- ☐ Chapter 2
- ☐ Chapter 3

## Summary

# Meet the Team



Edward Vincent **"VINCE"** Duero

Fave Childhood Series: **Naruto**  
Fave Series Now: **Silicon Valley**



Rosiel Jazmine **"ROSE"** Villareal

Fave Childhood Series: **Avatar - The Last Airbender**  
Fave Series Now: **The Good Place**



Jericho Carlo **"ECHO"** Agudo

Fave Childhood Series: **Detective Conan**  
Fave Series Now: **Bridgerton**

# SageMaker Deployment Options

Real-time Inference Endpoint  
Batch Transform

# Real-Time Inference Endpoint

**curl -X POST -H**

```
"Content-Type:application/json;  
format=pandas-split"  
--data '{"columns":["alcohol",  
"chlorides",...],"data":[[12.8,  
0.029,...]]}' http://<endpoint>
```

Request

**Real-time  
Inference  
Endpoint**

Response

**[6.379428821398614]**

# Use SageMaker Real-Time Inference

- For inference workloads when you have **real-time, interactive, low latency** requirements
  - For providing a **persistent, real-time endpoint**
- When you want to **call an endpoint** that accepts data input and outputs a model prediction in real-time, i.e., it **gives you a response anytime you want**

# SageMaker Real-time Inference Hosting Options



**Single** model



**Multiple** models in **one container** behind one endpoint



Multiple models for **serial inference pipeline** behind one endpoint



**Multiple** models w/c use **different containers** behind one endpoint

# SageMaker Real-time Inference Features

Auto **Scale** Models

Host **Storage** Volumes (EBS)

Test Model **Versions** using Production Variants

Note: Specify traffic distribution or variants

**Monitor** Drifts in Data, Model, Bias, Feature Attribution

Note: Currently supports only tabular data, single model endpoint



# SageMaker Real-time Inference Use Cases

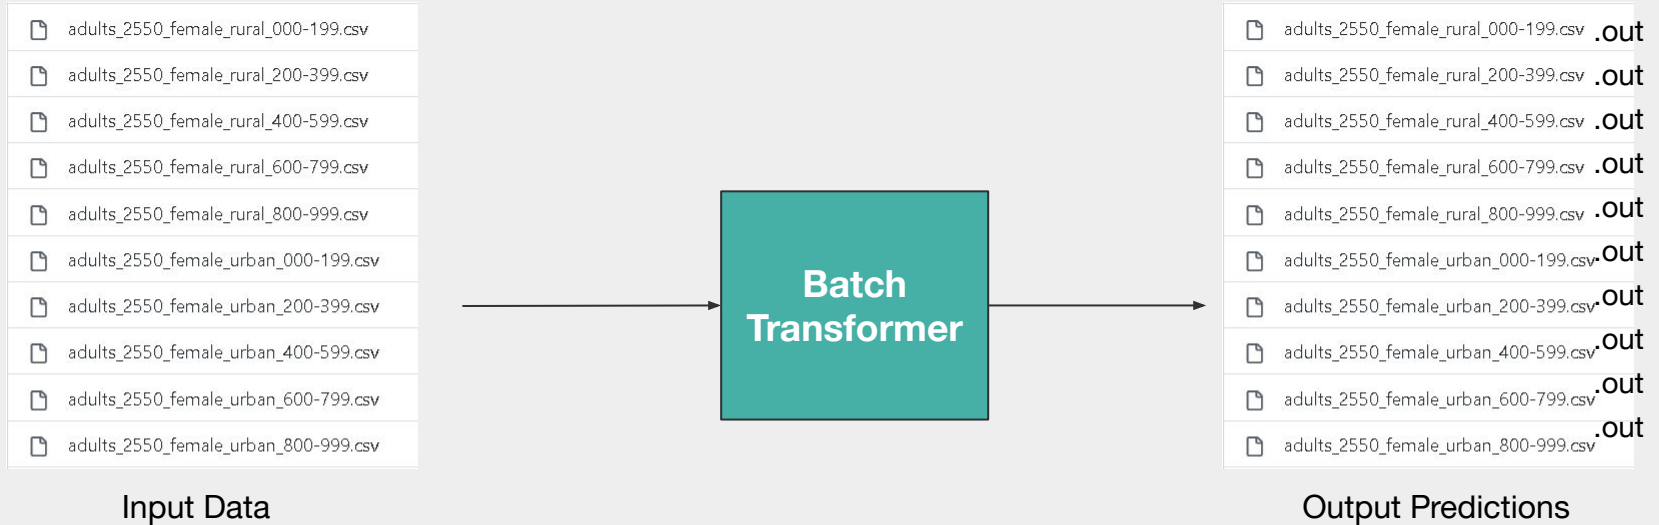
## Real-Time Plant Disease Detection



## Real-time Fraud Detection



# Batch Transform



# Use Batch Transform

- For getting inferences for **large datasets**
- For **preprocessing** datasets to remove noise or bias
- For running inference when a **persistent endpoint is not needed**
- For associating input records with inferences to assist **results interpretation**

- For getting **predictions** for datasets in **bulk or batches** instead of streams
- For **preprocessing entire datasets**

# How SageMaker Batch Transform Works

**Distributes workload** between compute instances

Processes **one file per instance** or splits a **file** into mini-batches

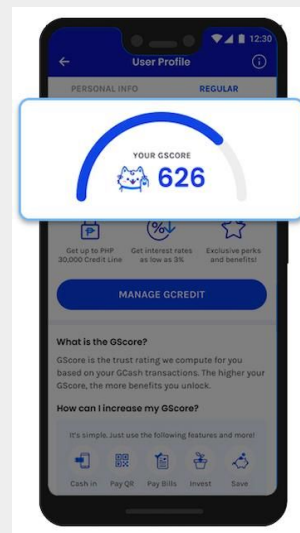
Stores **output files** in a specified location, e.g. S3 bucket, with same names as input files + `.out` file extension

# SageMaker Batch Transform Use Cases

## NLP Preprocessing

x	blog	eshop	faq	fpage
blog	24	1	0	1
eshop	1	6	1	2
faq	0	1	3	0
fpage	6	2	2	4
listing	3	2	3	1
php	2	0	0	4
spage	1	3	0	3

## Credit Scoring



## Real-time Inference

- Persistent - instances stay running until shut down
- Predicts for individual records in real-time

## Batch Transform

- Instances torn down when job completes
- Predicts for groups of records

# Amazon SageMaker Cookbook

# Chapter 1 Summary

- gentle introduction in using Amazon SageMaker

## *What we learned:*

- How to perform a basic machine learning project in Amazon SageMaker
  - How to perform a machine learning project the AWS way
- How to deploy an inference endpoint for the trained model and how to invoke the endpoint to retrieve predictions



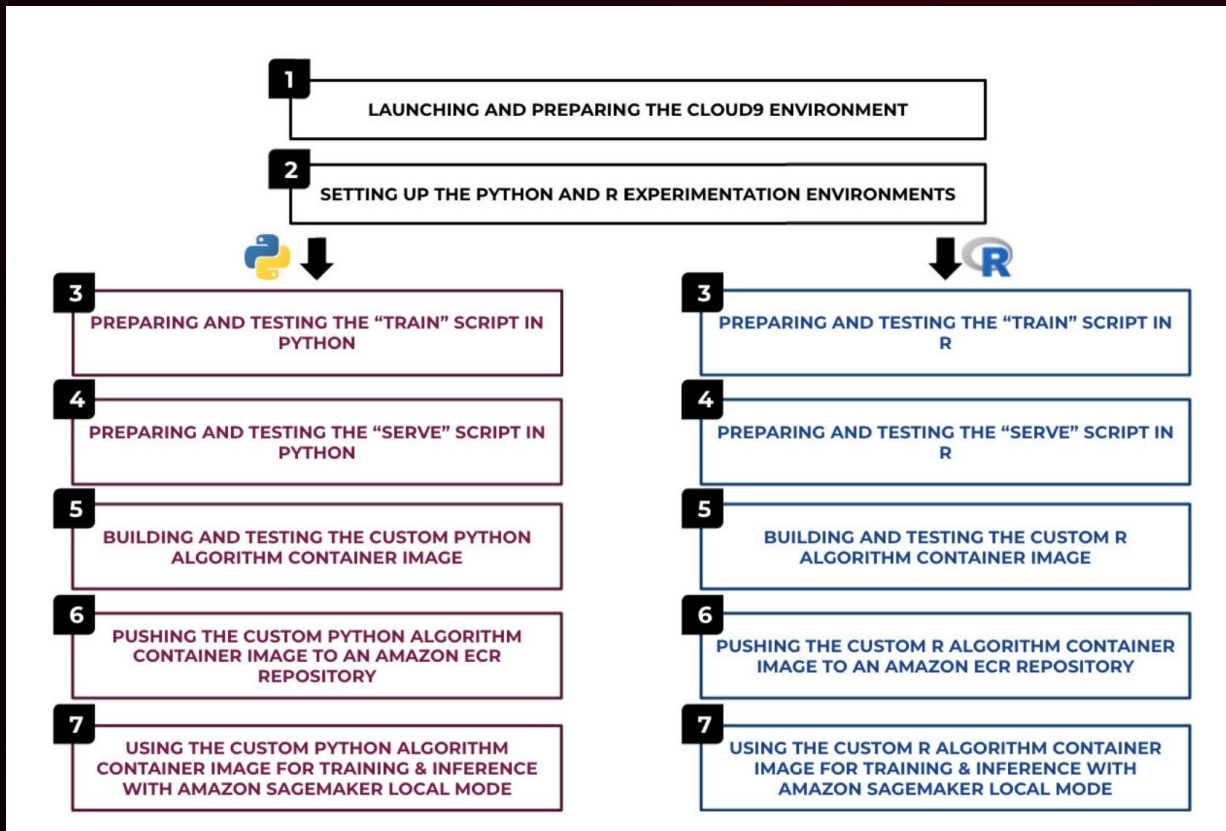
# Chapter 2 Summary

- deep dive into how Amazon SageMaker works internally

## *What we learned:*

- How to implement a custom model in Python and R
  - Create train and serve scripts
  - Build docker image from scripts and upload to a repository
  - Use image to create a machine learning model in a project

# Chapter 2 Summary



# Chapter 3 Summary

- demonstration on the different machine learning libraries and frameworks in Amazon Sagemaker


## *What we learned:*

- How to generate a synthetic dataset
- How to train and deploy a model using the following:
  - TensorFlow and Keras
  - PyTorch
  - scikit-learn

# Chapter 3 Summary - Synthetic dataset

```
def formula(x)
    if x >= -2000:
        return x
    else:
        return -x - 4000
```

```
def generate_synthetic_data(n_samples=1000,
                             start=-5000, end=5000):
    np.random.seed(42)
    x = np.random.randint(
        low=start,
        high=end,
        size=(n_samples,)).astype(int)
    y = np.vectorize(formula)(x) + \
        np.random.normal(150, 150, n_samples)
    return (x,y)
```



# Chapter 3 Summary - Libraries and frameworks

## TensorFlow and Keras

### Entrypoint:

- can use the environment variables set by SageMaker for the entrypoint script.

### Deployment:

- can directly use `deploy()` function.

### Notes:

- industry-focused

## PyTorch

### Entrypoint:

- separate inference script for deployment.

### Deployment:

- PyTorchModel needs to be used and initialized.

### Notes:

- research-oriented

## scikit-learn

### Entrypoint:

- `model_fn()` function needs to be defined.

### Deployment:

- can directly use `deploy()` function.

### Notes:

- simple and beginner friendly

# Summary

## SageMaker Deployment Options

- ✓ Real-Time Inference Endpoint
- ✓ Batch Transform

## Amazon SageMaker Cookbook

- ✓ Chapter 1
- ✓ Chapter 2
- ✓ Chapter 3

**apper.ph**